

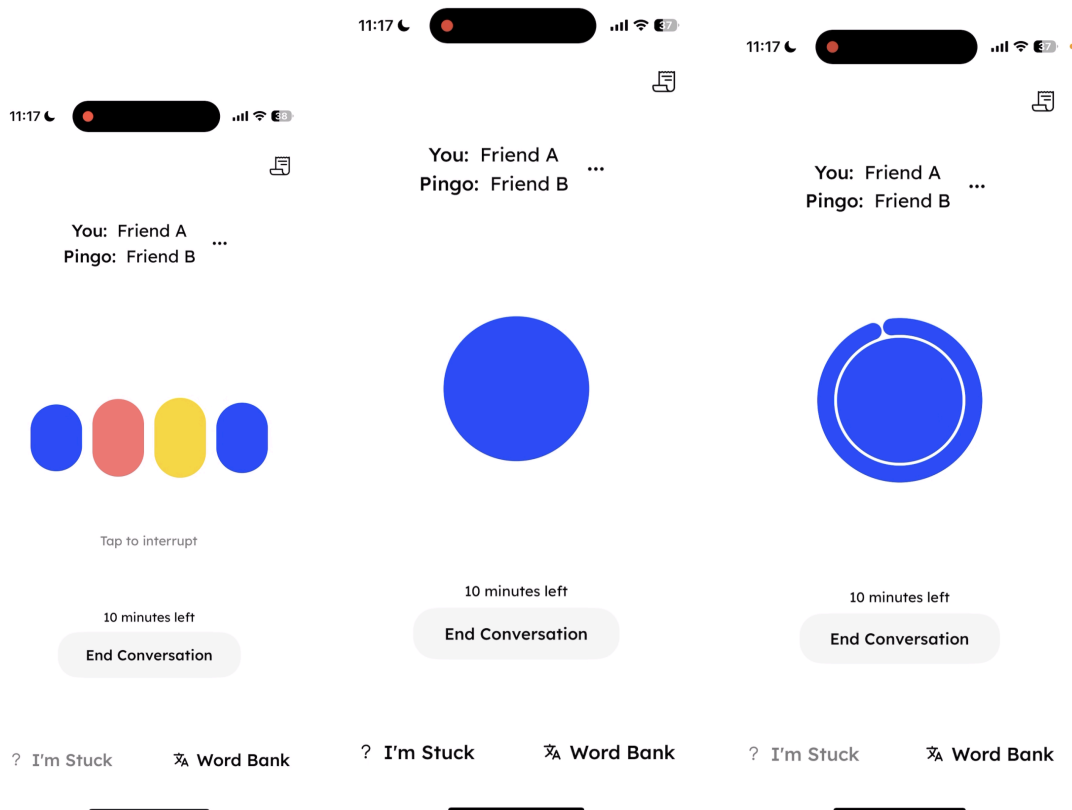
Two Part Take Home Project: Mini Version of Pingo AI w/ OpenAI Realtime API Integration

Short Description: Build a basic web version of Pingo AI. The user should be able to:

- Select a scenario before starting.
- Speak to an AI in real time.
- See audio frequency visualizations while the AI responds.

You'll build both the frontend and a minimal backend that connects directly to the OpenAI Realtime API.

Design Assets: Below are screenshots of the Pingo app that can be used for design. The Pingo colors are [#1F4DFE](#), [#FF6F6F](#), and [#FFD400](#). You can skip the "hold-to-record" animation if it's not feasible.



Note: You can use <https://youtu.be/6CP9Iyyxsqg> for a visualization of how interaction works

Success Criteria:

- User can choose the scenario before speaking.
- Hold to speak.
- Stream audio from OpenAI.

- Audio visualization (frequency bars) reflects live audio data.
- Basic speaking animations tied to playback.
- Smooth, gap-free audio playback.
- Clear loading and interaction states.
- Simple, intuitive UX.

Stretch Goals (Any are totally optional):

- Show the conversation transcript.
- Provide a summary and feedback of the conversation at the end.
- Allow choosing different languages.
- Any tool calling feature you want – e.g., AI ends the conversation when told.

Stack:

- React JS + Tailwind CSS or just React styling
- Node JS/Express